

Minesweeper Solver Analysis

Jacob Slagle

January 25, 2019

Introduction

Minesweeper Solver is a personal project I started to practice my coding skills. Specifically I wanted to take a problem, develop an algorithm to solve it, and create a project around it with all of the trappings. The game minesweeper was suitably complex: a good, efficient algorithm would take a bit of thinking, but not so much that I couldn't also focus on code design and increasing my Python knowledge. I also like the poetry of taking something which I've used as a distraction in the past and using it to improve myself.

The Problem

Simply put, the problem is to solve a game of Minesweeper (I assume the reader is familiar with the mechanics of the game). In more detail, I seek an algorithm that, if possible, wins a game of Minesweeper without ever risking losing the game, i.e. never makes a move unless it's a sure move. It suffices to look at a freeze frame of the game and determine with certainty which spaces must contain mines, and so should be flagged (right-clicked), and which spaces must not contain mines, and so should be revealed (left-clicked). It's unnecessary to find all such spaces, because an algorithm that can find at least one new space each time can be applied repeatedly.

What if there are no spaces that can be determined to be mined or free of mines (henceforth "free" for short)? Indeed, this is the case at the beginning of the game before any spaces have been revealed. However every game of minesweeper I've played, including the one I've coded, ensures that the first square selected is free as well as all surrounding squares. As it turns out, sometimes it's impossible to make any certain moves in the middle of the game either. In this case all we can do is guess, and I will discuss below how to make a good guess. For now though, we consider this: when possible, how do I determine which moves I can make with certainty.

Theory and Algorithms

Here I use the pronoun "we" as in formal mathematics.

As a starting point we consider what exactly it means to beat a game of minesweeper. To win we must place flags on exactly those squares which contain mines and reveal the rest of the squares- in other words we must find the right "placement" of flags on the board. To be slightly more formal, a **placement** is a subset squares on the board which are neither revealed nor flagged where we are considering placing flags (perhaps in addition to the flags already on the board). Now suppose we are playing a particular game and furthermore we are looking at a freeze frame of the game where some of the squares are revealed and perhaps some squares are flagged. The revealed squares in this frame contain numbers which tell us how many mines are adjacent to it. Since we aim to place flags on mines and only mines, these number hints are constraints on the way flags can be placed (we already placed on the board are correctly placed). With an eye towards satisfying these constraints, we make the following definition:

Definition 1. A placement is **satisfactory** if after placing all flags in it we have this property: for every revealed square s , the number hint h in s is equal to the number of flags surrounding s .

Observe that the correct placement of flags on mines is itself a satisfactory placement. It follows that if a square is flagged in *every* satisfactory placement, it must contain a mine. Therefore we say placing a flag on such a square is a **certain flag** (a flag placed with certainty). Likewise, if a square is *unflagged* in every mine placement, we know it must be free (i.e. it does not contain a mine). Thus we say revealing is a **certain reveal**. A **certain move** is a certain flag or a certain reveal. Consider the value of a certian move over an uncertain move: in playing Minesweeper, we do not know where the mines are but we can figure out what the satisfactory placements by looking at number hints, and so can deduce certain moves. Moreover we really *have* to resort to looking only at certain moves because if a reveal is uncertain, there are satisfactory placements in which the revealed square contains a mine, and if one of those satisfactory placements is the actual mine placement then revealing the square loses the game. Uncertain flags, while not immediately fatal, can lead to errors down the line. Hence we only want to make certain moves.

This line of reasoning lends itself to an algorithm idea: generate all satisfactory placements, check which squares are consistently flagged or unflagged across each placement and make certain moves accordingly. Let's assume for the moment we have a function `satisfactoryPlacements(G)` which takes a game of minesweeper G and generates all satisfactory placements.

```

function SOLVE( $G$ )                                     ▷ Version 1
   $Mines \leftarrow$  game board
   $Free \leftarrow$  game board
  for  $P$  in satisfactoryPlacements( $G$ ) do
     $Mines \leftarrow Mines \cap P$ 
     $Free \leftarrow Free - P$ 
  end for
  flagAll( $Mines$ )
  revealAll( $Free$ )

```

end function

We just need to implement `satisfactoryPlacements()`. We start with a brute force approach: go over all mine placements and yield the ones that are satisfactory. Mine placements are just subsets of the game board B , i.e. elements of the powerset $\mathcal{P}(B)$.

```

function SATISFACTORYPLACEMENTS( $G$ )                                ▷ Version 1
  for  $P$  in  $\mathcal{P}(B)$  do
    if  $P$  is satisfactory then
      yield  $P$ 
    end if
  end for
end function

```

This method checks $|\mathcal{P}(B)| = 2^b$ placements, where $b = |B|$ is the size of the board. We can check if a placement is satisfactory in linear time, and since the placements are bound by the board in size, we have that `satisfactoryPlacements()` has an $O(b2^b)$ runtime. The other parts of `solve()`. The other work in `solve` is relatively insignificant and so this also the runtime for `solve`. If you've ever played minesweeper though it might occur to you that it's unnecessary to look at the entire board, since in a sense we only have information about the unrevealed squares which are adjacent to numbered squares. With this in mind we make the following definitions: a square is **in-play** if (in the frame in question) it is unrevealed and unflagged and adjacent to a revealed square. We will refer to all in-play squares collectively as the **perimeter** because they circumscribe the revealed part of the board.

Lemma. *If S is a satisfactory placement and P is the perimeter, then $S' = S \cap P$ is also a satisfactory placement.*

Proof. Let r be a revealed square. All the flags in S adjacent to r are also in $S' = P \cap S$ because the perimeter P contains all the in-play squares in the current frame. In particular r has the same number of flags adjacent to it in P' as in P . This goes for each revealed square r , so S' like S is a satisfactory placement. \square

Claim 1. 1. *A flag is certain if and only if it occurs in each satisfactory placement contained within the perimeter. Additionally, certain flags always occur within the perimeter*

2. *A reveal is certain if and only if it is in the perimeter and the space to be revealed is unflagged in each satisfactory placement contained within the perimeter.*

3. *Certain moves occur only in the perimeter of the board.*

Proof. 1. If a flag is certain if it is contained in every satisfactory placement, in particular it is a part of the satisfactory placements contained within the perimeter. Conversely, suppose s is a space that is flagged in every

satisfactory placement that is contained within the perimeter. Let S be an arbitrary satisfactory placement. Set S' be the intersection of S and the perimeter. By the lemma, S' is also satisfactory, and so by our hypothesis, $s \in S'$. Of course $S' \subset S$, and so $s \in S$. Thus s is in every satisfactory placement S , meaning s is a certain flag. This proves the iff statement.

To prove that all certain flags are in the perimeter let M be the actual mine placement, which is satisfactory (for the purposes of this proof M could be any satisfactory mine placement). By the lemma M' , the intersection of M and the perimeter is satisfactory. By the definition of certain, all certain flags are in M' which is in the perimeter.

2. Suppose s is a certain reveal. By definition then, s is not contained in any satisfactory mine placement, in particular those contained within the perimeter. Suppose for contradiction s was not in the perimeter. Let M be the actual mine placement (or any satisfactory mine placement for that matter), and set $M' = M \cup \{s\}$. Since s is not in the perimeter it is not in play meaning that adding it to a placement would not change the number of flags around any revealed square. In particular adding it to M would yield a placement $M' = M \cup \{s\}$ which had the same flag counts around each revealed square, i.e. another satisfactory placement. Thus s is in a satisfactory placement M' , which is a contradiction because s is a certain reveal.

Conversely, suppose s is in the perimeter and is unflagged in each satisfactory placement contained within the perimeter. Let S be a satisfactory placement, and let S' be its intersection with the perimeter. By the lemma, S' is a satisfactory placement, and so by our hypothesis $s \notin S'$. But since s is in the perimeter and S' is the intersection of S with the perimeter, the only way $s \notin S'$ is if $s \notin S$. This shows that s is contained in no satisfactory placement and so by definition s is a certain reveal.

3. This statement follows immediately from the first two. □

This claim shows us that our algorithm need only focus on the perimeter as opposed to the whole board. That is, we can initialize mines and free to be just the perimeter and we only need to generate satisfactory placements that are contained within the perimeter. Here are the revised versions of `solve()` and `satisfactoryPlacements()`:

```

function SOLVE( $G$ ) ▷ Version 2
   $Mines \leftarrow \text{perimeter}(G)$ 
   $Free \leftarrow \text{perimeter}(G)$ 
  for  $P$  in satisfactoryPlacements( $G$ ) do
     $Mines \leftarrow Mines \cap P$ 
     $Free \leftarrow Free - P$ 
  end for
  flagAll( $Mines$ )

```

```

    revealAll(Free)
end function
function SATISFACTORYPLACEMENTS(G) ▷ Version 2
    for P in  $\mathcal{P}(\text{perimeter}(G))$  do
        if P is satisfactory then
            yield P
        end if
    end for
end function

```

This reduces the runtime to $O(b2^p)$ where b is the size of the board as before and p is the size of the perimeter. While this is still exponential, the perimeter is *much* smaller than the board except in pathological cases. In practice though (e.g. when using the algorithms outlined in this document), the revealed part of the board is a 2-dimensional contiguous "blob" and so gives rise to a 1-dimensional perimeter. This makes a huge difference given that it acts as an exponent in the runtime.

It still seems like unnecessary work is being done. Even when we restrict our attention to the perimeter, it seems like we're doing extra work when we go back and check if a placement is satisfactory- after all this is where the linear factor b in $O(b2^p)$ comes from. What if rather than looking at every placement, we "built up" satisfactory placements according to the hints on the board. If we do it correctly we don't have to check that it we are getting exactly the satisfactory placements. We introduce another term: the **fringe** consists of all revealed squares (numbers) with at least one adjacent square that is neither revealed nor flagged. Note the distinction between the fringe and the perimeter: the fringe consists of the edge (or edges), of the revealed part(s) of the board, while the perimeter is unrevealed and wraps around the fringe. Functionally speaking the fringe contains all the hints on the board that have yet to be satisfied, and as such is our focus if we are building up satisfactory placements.

The core of the algorithm is to look at a point in the fringe and add the appropriate number of flags around it to satisfy it. There are usually multiple ways to add the appropriate number of flags, so we choose one way and proceed to the next point, doing the same with this point. We continue in this manner until we have a satisfactory placement, which we then yield. We then backtrack and consider other ways of adding flags around points. We now present the pseudocode.

```

function SATISFACTORYPLACEMENTS(G) ▷ Version 3
    SPHelper(G, 0,  $\emptyset$ )
end function
function SPHELPER(G, i, placement)
    if  $i = \text{length}(G.\text{fringe})$  then ▷ recursion base case
        yield placement
    end if
    point  $\leftarrow G.\text{fringe}[i]$  ▷ Get the ith point in the fringe
    if point has too many flags around it when including those in placement

```

```

then
    return ▷ Go back up one level of the call stack
else if point already has the right number of flags around it then
    SPHelper( $G, i + 1, placement$ )
else
    for all sets of flags newFlags that can be added to satisfy point with-
    out adding flags around  $G.fringe[0], G.fringe[1], \dots, G.fringe[i - 1]$  do
         $placement \leftarrow placement \cup newFlags$ 
        SPHelper( $G, i + 1, placement$ )
         $placement \leftarrow placement - newFlags$ 
    end for
end if
end function

```

Claim 2. *The above version of SatisfactoryPlacements() yields exactly those satisfactory placements contained within the perimeter.*

Proof. The perimeter is just the union of all in-play squares adjacent to points in the fringe, so we might rephrase the claim as SatisfactoryPlacements() yields exactly those satisfactory placements that contain only flags adjacent to points in the fringe. Now consider the following propositions which depends on the variable k .

OnlySatisfactory(k): Whenever a call of the form SPHelper($G, k, placement$) is made, which is to say a call with $i = k$, then the variable *placement* is a placement that satisfies the points $G.fringe[0], \dots, G.fringe[k - 1]$

AllSatisfactory(k): Each placement which (a) satisfies the points $G.fringe[0], \dots, G.fringe[i - 1]$ and (b) contains only flags adjacent to $G.fringe[0], \dots, G.fringe[i - 1]$ is stored in the variable *placement* at the beginning of one of the function calls of the form SPHelper($G, k, placement$)

The statement OnlySatisfactory(length($G.fringe$)) should be interpreted as follows: when the base case of the recursion is entered (because i equals length($G.fringe$)) the *placement* variable contains a satisfactory placement because all points in $G.fringe$ are satisfied, and thus the placement that's yielded is satisfactory. The statement AllSatisfactory(length($G.fringe$)) guarantees that all satisfactory placements contained within the perimeter are yielded. Together these statements tell us that SatisfactoryPlacements() yields exactly those satisfactory placements contained within the perimeter. We prove them by induction.

There are no points to be satisfied in base case ($k = 0$), so OnlySatisfactory(0) is trivial. recursive calls are made with strictly ascending values of i , so SPHelper() is only called once with $i = 0$ once and in this case $placement = \emptyset$. Indeed \emptyset is the only placement which satisfies conditions (a) and (b) in AllSatisfactory(0)

Now suppose $\text{OnlySatisfactory}(k)$ is true. Calls of the form $\text{SPHelper}(G, k+1, \text{placement})$ are only made within calls of the form $\text{SPHelper}(G, k, \text{placement})$. When such a call is made from the else if clause, then placement already satisfies $G.\text{fringe}[k]$ and $\text{OnlySatisfactory}(k)$ tells us that it also satisfies $G.\text{fringe}[0], \dots, G.\text{fringe}[k-1]$, and together these give us $\text{OnlySatisfactory}(k+1)$. Otherwise the call $\text{SPHelper}(G, k+1, \text{placement})$ is made within the for each loop. Adding newFlags to placement guarantees that placement satisfies $G.\text{fringe}[k]$, but by the condition in the for each loop newFlags doesn't contain any flags around $G.\text{fringe}[0], \dots, G.\text{fringe}[k-1]$. $\text{OnlySatisfactory}(k)$ guarantees that placement satisfied $G.\text{fringe}[0], \dots, G.\text{fringe}[k-1]$ to begin with, and after adding newFlags which doesn't contain any flags around these points, they continue to be satisfied. Thus $\text{OnlySatisfactory}(k+1)$ is true.

Now suppose $\text{AllSatisfactory}(k)$ is true. Let P be a placement as described in $\text{AllSatisfactory}(k+1)$. Let N be the flags from P which are adjacent to $G.\text{fringe}[k]$ but are not adjacent to $G.\text{fringe}[0], \dots, G.\text{fringe}[k-1]$ (N may be empty), and set P' equal to $P - N$. It should be clear that P' is one of the placements described in $\text{AllSatisfactory}(k)$, and so there is a call of the form $\text{SPHelper}(G, k, \text{placement})$ with placement equal to P' . If $N = \emptyset$ then $\text{placement} = P' = P$ satisfies $G.\text{fringe}[0], \dots, G.\text{fringe}[k]$ already and the else if clause makes the call $\text{SPHelper}(G, k+1, \text{placement})$ as required. Otherwise if $N \neq \emptyset$, one of the iterations of the for each loop sets newMines equal to N and adds newMines to placement . Since placement equals P' before this it now equals $P' \cup N = P$. Thus we call $\text{SPHelper}(G, k+1, \text{placement})$ with $\text{placement} = P$ as desired. \square

Now to analyze the runtime. First observe that each iteration of the for each loop includes a recursive call to SPHelper , so we can associate the work done in this iteration as a part of the recursive call (i.e. we associate the work with the call $\text{SPHelper}(G, i+1, \text{placement})$, not with the current call). All other work done in call to SPHelper is constant, thus the runtime is proportional to the number of calls to SPHelper . At the beginning of each call to SPHelper placement stores a different subset of the perimeter (we could prove this by induction), and so the number of calls is bounded by the number of subsets of the perimeter, i.e. 2^p where p is the size of the perimeter. Thus the runtime is $O(2^p)$, and as desired we've dropped the factor of p from the previous version of $\text{SatisfactoryPlacements}()$. Also note that while the previous version of $\text{SatisfactoryPlacements}()$ tried literally every subset of the perimeter, this version cuts out a lot of subsets by not building upon placements as soon as it's clear that won't satisfy a point in the fringe, viz "if(*point* has too many flags around it)."